

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
In [6]: df=pd.read_csv("uber-raw-data-sep14.csv")
df.head()
```

```
Out[6]:
```

	Date/Time	Lat	Lon	Base
0	9/1/2014 0:01:00	40.2201	-74.0021	B02512
1	9/1/2014 0:01:00	40.7500	-74.0027	B02512
2	9/1/2014 0:03:00	40.7559	-73.9864	B02512
3	9/1/2014 0:06:00	40.7450	-73.9889	B02512
4	9/1/2014 0:11:00	40.8145	-73.9444	B02512

```
In [7]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1028136 entries, 0 to 1028135
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Date/Time    1028136 non-null object
1   Lat          1028136 non-null float64
2   Lon          1028136 non-null float64
3   Base         1028136 non-null object
dtypes: float64(2), object(2)
memory usage: 31.4+ MB
```

```
In [8]: df.isna().sum()
```

```
Out[8]:
```


<b>Date/Time</b>	0
<b>Lat</b>	0
<b>Lon</b>	0
<b>Base</b>	0

**dtype:** int64

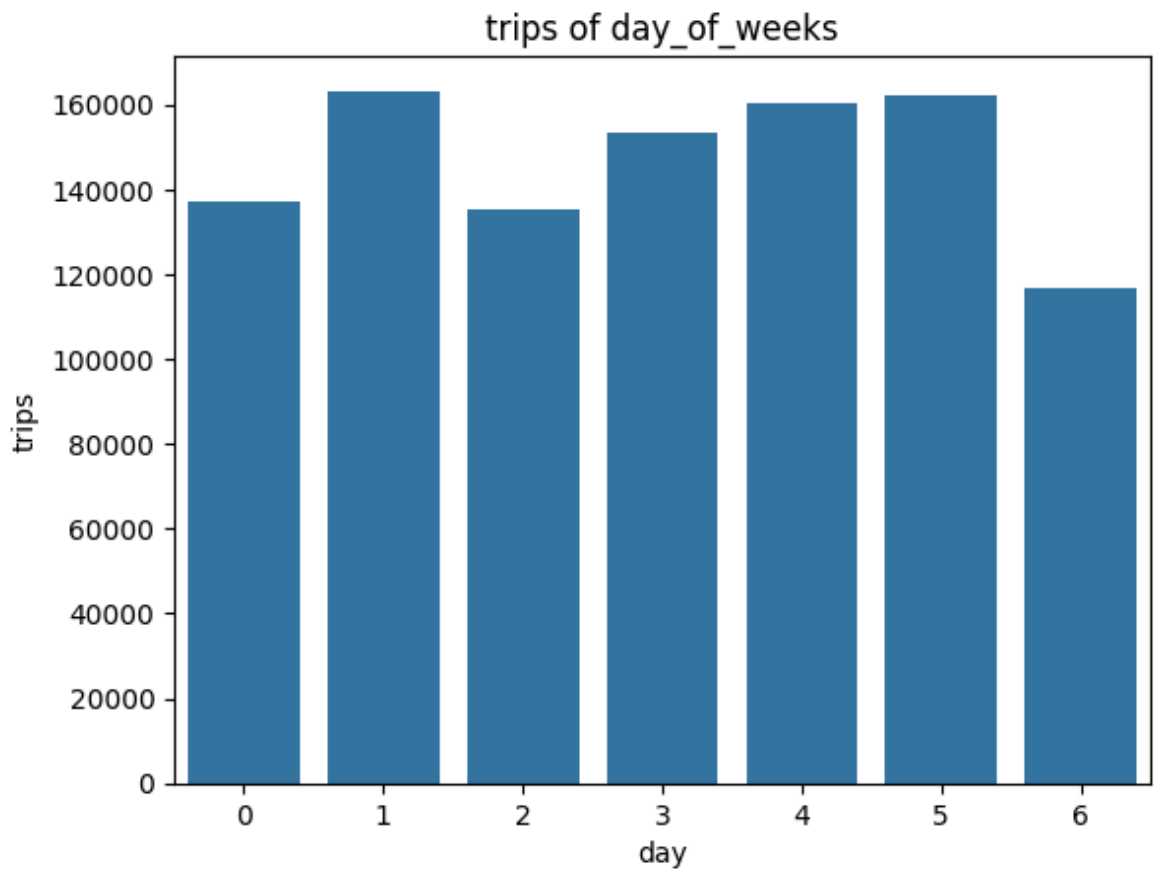
```
In [9]: df['Date/Time']=pd.to_datetime(df['Date/Time'],format='%m/%d/%Y %H:%M:%S')
```

```
In [10]: df['hour']=df['Date/Time'].dt.hour
df['day_of_week']=df['Date/Time'].dt.day_of_week
df['day_of_month']= df['Date/Time'].dt.day
df['month']=df['Date/Time'].dt.month
```


```
In [69]: sns.countplot(x='hour', data=df)
plt.title("trips Per hour")
plt.xlabel("hour")
plt.ylabel("trips")
plt.show()
```



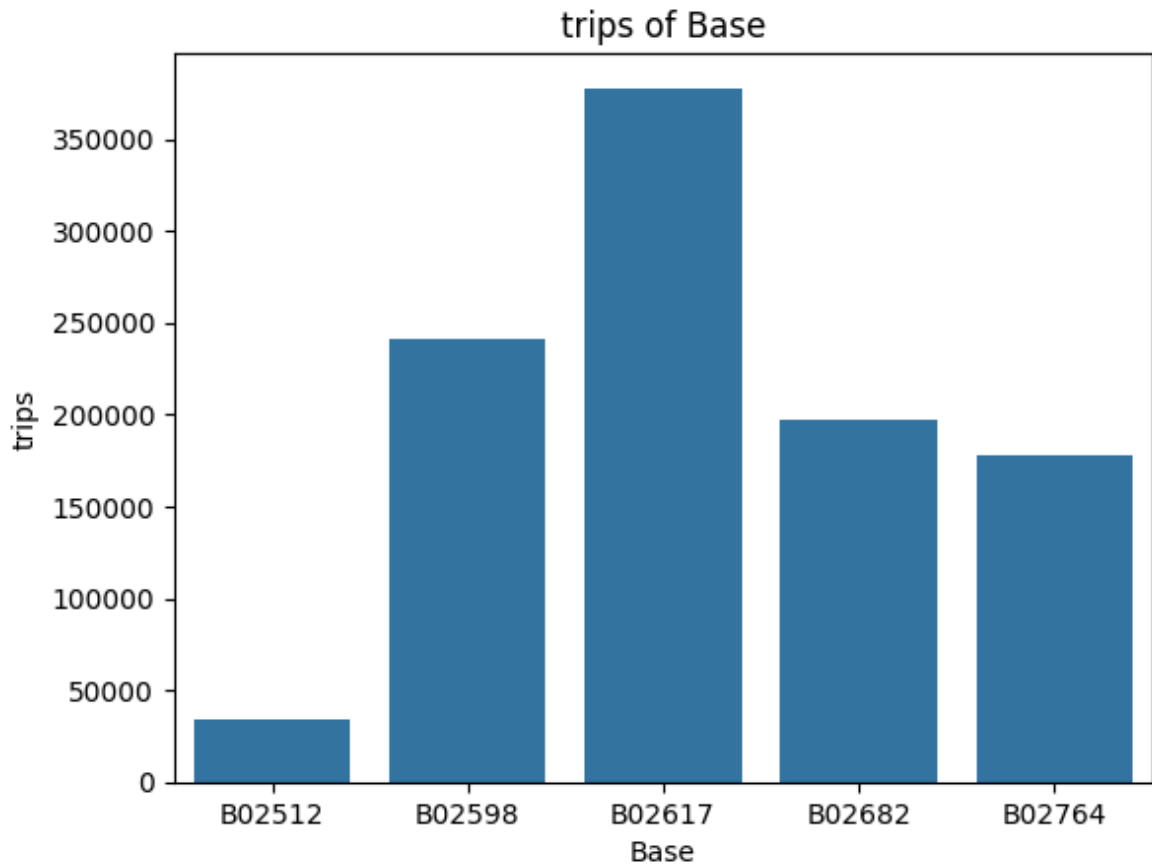
```
In [12]: sns.countplot(x='day_of_week', data=df)
plt.title("trips of day_of_weeks")
plt.xlabel("day")
plt.ylabel("trips")
plt.show()
```



```
In [13]: sns.countplot(x='day_of_month', data=df)
plt.title("trips of day_of_months")
plt.xlabel("day")
plt.ylabel("trips")
plt.show()
```



```
In [14]: sns.countplot(x='Base',data=df)
plt.title("trips of Base")
plt.xlabel("Base")
plt.ylabel("trips")
plt.show()
```



```
In [17]: df.set_index('Date/Time', inplace=True)
trips_per_hour = df.resample('H').size().reset_index(name='trip_count_per_hour')
trips_per_hour.name = 'Trip Count/hour'
trips_per_day = df.resample('D').size()
trips_per_day.name = 'Trip Count/day'
trips_per_month = df.resample('M').size()
trips_per_month.name = 'Trip Count/month'
```

```
In [18]: from sklearn.preprocessing import LabelEncoder

le=LabelEncoder()
df['Base'] = le.fit_transform(df['Base'])
```

```
In [78]: df['is_weekend'] = df['day_of_week'].isin([5, 6]).astype(int)
df['is_night_hour'] = df['hour'].isin(list(range(0, 6)) + list(range(22, 24))).a

df_hourly = df.resample('H').agg({
    'Lat': 'mean',
    'Lon': 'mean',
    'Base': lambda x: x.mode()[0] if not x.mode().empty else np.nan,
    'hour': 'first',
    'day_of_week': 'first',
    'day_of_month': 'first',
    'month': 'first',
    'is_weekend': 'first',
```

```

    'is_night_hour': 'first'
}).reset_index()

trip_counts = df.resample('H').size().reset_index(name='trip_count_per_hour')
df_model = pd.merge(df_hourly, trip_counts, on='Date/Time')

df_model['lag_1'] = df_model['trip_count_per_hour'].shift(1)
df_model['rolling_mean_3h'] = df_model['trip_count_per_hour'].rolling(3).mean()
df_model['rolling_std_6h'] = df_model['trip_count_per_hour'].rolling(6).std()
df_model['hour_sin'] = np.sin(2 * np.pi * df_model['hour'] / 24)
df_model['hour_cos'] = np.cos(2 * np.pi * df_model['hour'] / 24)

df_model.dropna(inplace=True)

```

```

In [79]: df_model['lag_2']=df_model['trip_count_per_hour'].shift(2)
df_model['lag_3']=df_model['trip_count_per_hour'].shift(3)
df_model['rolling_mean_6h']=df_model['trip_count_per_hour'].rolling(6).mean()
df_model['rolling_mean_12h']=df_model['trip_count_per_hour'].rolling(12).mean()
df_model.dropna(inplace=True)

X_new = df_model[['lag_1','lag_2','lag_3','rolling_mean_3h','rolling_mean_6h','i
y_new = df_model['trip_count_per_hour']
train_size=int(len(X_new)*0.8)
X_train_new,X_test_new=X_new[:train_size],X_new[train_size:]
y_train_new,y_test_new=y_new[:train_size],y_new[train_size:]

```

```

In [56]: from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error,r2_score,mean_absolute_error
rf=RandomForestRegressor()
rf.fit(X_train_new,y_train_new)
pred=rf.predict(X_test_new)

mse=mean_squared_error(y_test_new,pred)
mae=mean_absolute_error(y_test_new,pred)
r2=r2_score(y_test_new,pred)
print("MSE:",mse)
print("MAE:",mae)
print("R2:",r2)

```

MSE: 7774.557685227275


MAE: 66.59625

R2: 0.9848226195942449

```

In [57]: plt.plot(y_test_new.values,label='test')
plt.plot(pred,label='pred')
plt.legend()
plt.show()

```



```
In [58]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error

lr=LinearRegression()

lr.fit(X_train_new,y_train_new)
predlr=lr.predict(X_test_new)

mse=mean_squared_error(y_test_new,pred)
mae=mean_absolute_error(y_test_new,pred)
r2=r2_score(y_test_new,pred)
print("MSE:",mse)
print("MAE:",mae)
print("R2:",r2)
```

MSE: 7774.557685227275

MAE: 66.59625

R2: 0.9848226195942449

```
In [62]: from sklearn.svm import SVR
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error

svr=SVR()
svr.fit(X_train_new,y_train_new)
predsvr=svr.predict(X_test_new)

mse=mean_squared_error(y_test_new,pred)
mae=mean_absolute_error(y_test_new,pred)
r2=r2_score(y_test_new,pred)
print("MSE:",mse)
print("MAE:",mae)
print("R2:",r2)
```

MSE: 7774.557685227275  
MAE: 66.59625  
R2: 0.9848226195942449

```
In [63]: from xgboost import XGBRegressor
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error

xgb=XGBRegressor()


xgb.fit(X_train_new,y_train_new)
pred_xgb=xgb.predict(X_test_new)

mse=mean_squared_error(y_test_new,pred)
mae=mean_absolute_error(y_test_new,pred)
r2=r2_score(y_test_new,pred)
print("MSE:",mse)
print("MAE:",mae)
print("R2:",r2)
```


MSE: 7774.557685227275  
MAE: 66.59625  
R2: 0.9848226195942449

```
In [64]: importance=rf.feature_importances_
feature_name=X_new.columns

plt.barh(feature_name,importance)
plt.show()
```




```
In [32]: plt.figure(figsize=(12, 5))
plt.plot(df_model['Date/Time'], df_model['trip_count_per_hour'])
plt.xlabel("Date")
plt.ylabel("Trips Per Hour")
plt.title("Hourly Trip Counts Over Time")
plt.show()
```



```
In [33]: from pandas.plotting import autocorrelation_plot

autocorrelation_plot(df_model['trip_count_per_hour'])
plt.show()
```



```
In [34]: print("RandomForest Params:", rf.get_params())

RandomForest Params: {'bootstrap': True, 'ccp_alpha': 0.0, 'criterion': 'squared_error', 'max_depth': None, 'max_features': 1.0, 'max_leaf_nodes': None, 'max_samples': None, 'min_impurity_decrease': 0.0, 'min_samples_leaf': 1, 'min_samples_split': 2, 'min_weight_fraction_leaf': 0.0, 'monotonic_cst': None, 'n_estimators': 100, 'n_jobs': None, 'oob_score': False, 'random_state': None, 'verbose': 0, 'warm_start': False}
```

```
In [80]: from sklearn.model_selection import RandomizedSearchCV

rf_param_grid= {
```



```

    'n_estimators': [100,200,300,500],
    'max_depth' : [5,10,15,20, None],
    'min_samples_split' : [2,5,10],
    'min_samples_leaf' : [1,2,4],
    'max_features' : ['auto','sqrt','log2']
}

xgb_param_grid = {
    'n_estimators': [100, 200, 300],
    'learning_rate': [0.01, 0.05, 0.1],
    'max_depth': [3, 5, 7],
    'subsample': [0.7, 1.0],
    'colsample_bytree': [0.7, 1.0]
}

rf=RandomForestRegressor(random_state=42)

random_search=RandomizedSearchCV(estimator=rf,param_distributions=rf_param_grid,
                                cv=3,verbose=2,random_state=42,n_jobs=-1)

xgb_random = RandomizedSearchCV(XGBRegressor(random_state=42, objective='reg:squ
                                xgb_param_grid,n_iter=10,cv=3,scoring='r2',verbo
                                n_jobs=-1)
xgb_random.fit(X_train_new, y_train_new)
xgb_best = xgb_random.best_estimator_
xgb_pred = xgb_best.predict(X_test_new)
random_search.fit(X_train_new,y_train_new)
best_model=random_search.best_estimator_
rf_pred=best_model.predict(X_test_new)

mse = mean_squared_error(y_test_new, rf_pred)
mae = mean_absolute_error(y_test_new, rf_pred)
r2 = r2_score(y_test_new, rf_pred)
mse_xgb = mean_squared_error(y_test_new, xgb_pred)
mae_xgb = mean_absolute_error(y_test_new, xgb_pred)
r2_xgb = r2_score(y_test_new, xgb_pred)
print("Best Parameters:", random_search.best_params_)
print("MSE:", mse)
print("MAE:", mae)
print("R2:", r2)

print("Best Parameters fpor xgb:", xgb_random.best_params_)
print("MSE for xgb:", mse_xgb)
print("MAE for xgb:", mae_xgb)
print("R2 for xgb:", r2_xgb)

```

Fitting 3 folds for each of 10 candidates, totalling 30 fits

Fitting 3 folds for each of 20 candidates, totalling 60 fits

Best Parameters: {'n\_estimators': 200, 'min\_samples\_split': 2, 'min\_samples\_leaf': 1, 'max\_features': 'sqrt', 'max\_depth': None}

MSE: 11621.48273865248

MAE: 73.81950354609928

R2: 0.9795435972839298

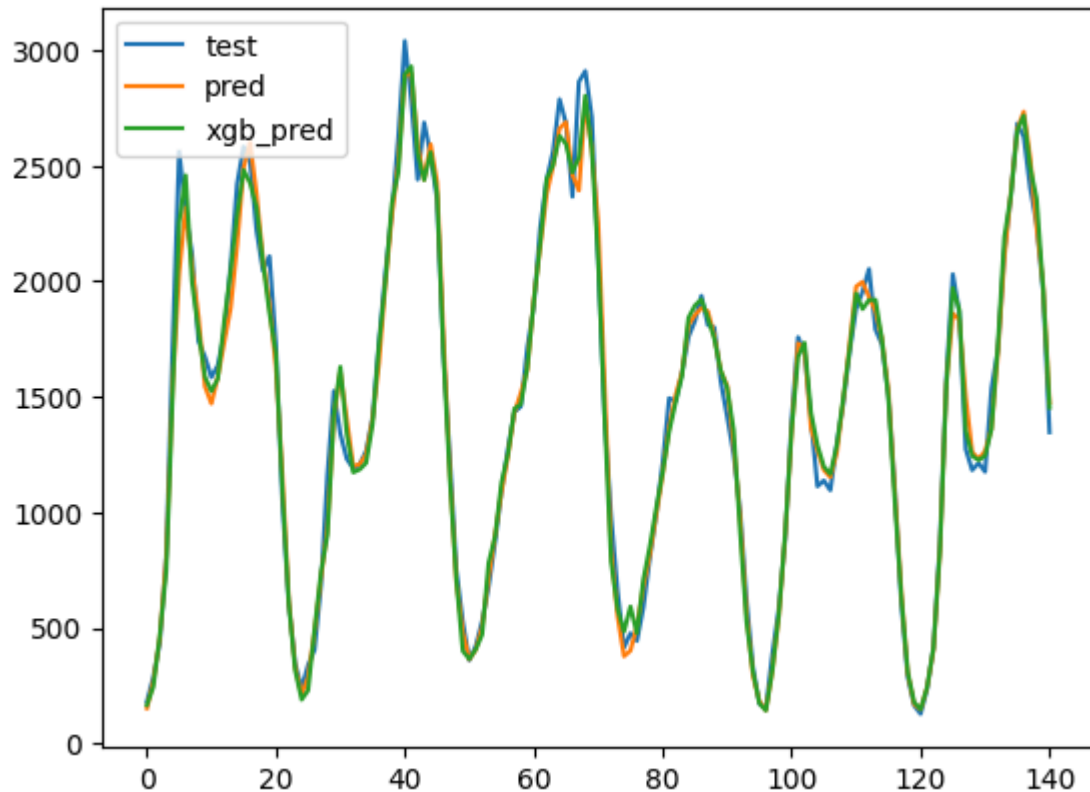
Best Parameters fpor xgb: {'subsample': 1.0, 'n\_estimators': 300, 'max\_depth': 3, 'learning\_rate': 0.1, 'colsample\_bytree': 1.0}

MSE for xgb: 9054.6396484375

MAE for xgb: 70.04219818115234

R2 for xgb: 0.9840618371963501

```
In [81]: plt.plot(y_test_new.values, label='test')
plt.plot(rf_pred, label='pred')
plt.plot(xgb_pred, label='xgb_pred')
plt.legend()
plt.show()
```



```
In [109... from sklearn.model_selection import RandomizedSearchCV
import random
from lightgbm import LGBMRegressor
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error
from scipy.stats import randint, uniform
gb_params = {
    'n_estimators': randint(100, 500),
    'learning_rate': uniform(0.01, 0.3),
    'max_depth': randint(3, 10),
    'min_samples_split': randint(2, 10),
    'min_samples_leaf': randint(1, 10),
    'subsample': uniform(0.6, 0.4)
}
lgb_params = {
    'n_estimators': randint(100, 500),
    'learning_rate': uniform(0.01, 0.3),
    'max_depth': randint(3, 10),
    'num_leaves': randint(20, 100),
    'subsample': uniform(0.6, 0.4),
    'colsample_bytree': uniform(0.6, 0.4)
}
gb_model = GradientBoostingRegressor()
lgb_model = LGBMRegressor()
gb_search = RandomizedSearchCV(
    gb_model, gb_params, n_iter=20, cv=3, scoring='neg_mean_squared_error', verbose=
)
gb_search.fit(X_train_new, y_train_new)
```

```

print("Best GB Parameters:", gb_search.best_params_)
gb_best = gb_search.best_estimator_
lgb_search = RandomizedSearchCV(
lgb_model, lgb_params, n_iter=20, cv=3, scoring='neg_mean_squared_error', verbose
)
lgb_search.fit(X_train_new, y_train_new)

print("Best LGB Parameters:", lgb_search.best_params_)
lgb_best = lgb_search.best_estimator_
# Initialize variables to store predictions for each model
gb_preds = []
lgb_preds = []

# Run RandomizedSearchCV for both models
for model_name, model, preds_list in [("Gradient Boosting", gb_best, gb_preds),
    preds = model.predict(X_test_new) # Make predictions
    preds_list.extend(preds) # Store predictions in the appropriate list

    # Evaluate performance (optional)
    mse = mean_squared_error(y_test_new, preds)
    mae = mean_absolute_error(y_test_new, preds)
    r2 = r2_score(y_test_new, preds)
    print(f"\n{model_name} Results:")
    print(f" MSE: {mse:.2f}")
    print(f" MAE: {mae:.2f}")
    print(f" R2 : {r2:.4f}")

```

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]



[illegible]

[illegible]

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf  
Best LGB Parameters: {'colsample\_bytree': np.float64(0.8473544037332349), 'learning\_rate': np.float64(0.12473859738014881), 'max\_depth': 6, 'n\_estimators': 419, 'num\_leaves': 22, 'subsample': np.float64(0.9439761626945282)}

Gradient Boosting Results:

MSE: 6431.33

MAE: 56.48

R<sup>2</sup> : 0.9887


LightGBM Results:

MSE: 8422.78

MAE: 63.48

R<sup>2</sup> : 0.9852

```
In [113... plt.plot(y_test_new.values,label='test')
plt.plot(lgb_preds,label='lgb_pred')
plt.plot(xgb_pred,label='xgb_pred')
plt.plot(gb_preds,label='gb_pred')
plt.plot(rf_pred,label='fr_pred')
plt.legend()
plt.show()
```



```
In [114... results = [
    {'Model': 'Gradient Boosting', 'MSE': 6431.33, 'MAE': 56.48, 'R²': 0.9887},
    {'Model': 'LightGBM', 'MSE': 8422.78, 'MAE': 63.48, 'R²': 0.9852},
    {'Model': 'Random Forest', 'MSE': 11621.48, 'MAE': 73.81, 'R²': 0.9795},
    {'Model': 'XGBoost', 'MSE': 9054.63, 'MAE': 70.04, 'R²': 0.9840}
]

results_df = pd.DataFrame(results)

results_df
```

Out[114...

	Model	MSE	MAE	R <sup>2</sup>
0	Gradient Boosting	6431.33	56.48	0.9887
1	LightGBM	8422.78	63.48	0.9852
2	Random Forest	11621.48	73.81	0.9795
3	XGBoost	9054.63	70.04	0.9840

In [ ]: